


Section 5: Web Concepts in System Design

Serialization: Data Exchange & Storage Formats


In this session we will explore serialization which is the fundamental process that enables applications, services and databases to exchange and store data efficiently across modern distributed systems.



In modern systems data is constantly moving between applications, services, databases and caches. The challenge is that computers work with in-memory objects but networks and storage systems need data in a transferable format; that gap is exactly what serialization solves. Serialization converts complex objects into a format that can be easily transmitted, stored and reconstructed elsewhere. Without serialization, communication between services, platforms, and programming languages would be difficult as well as inefficient. You encounter serialization almost everywhere in system design. APIs serialize data before sending responses, databases store serialized representation of objects, caching systems use it to persist data in-memory and distributed systems rely on it to exchange information across services. Now as systems grow and traffic increases, the choice of serialization format directly impacts performance, band width usage, storage efficiency and inter operability.

Understanding serialization is therefore much more than learning how to convert objects into bytes. It is about designing scalable, efficient, and reliable systems that can communicate seamlessly across services, platforms, and programming languages.

What is Serialization?



At its core serialization is about making data portable. Inside an application data exists as objects in-memory but memory structures cannot be sent directly across a network or stored in a database. Serialization transforms those objects into a standardized format that can be transmitted, persisted or shared with other systems.

When that data reaches its destination, deserialization performs the reverse operation, which is nothing but reconstructing the original object so that the receiving application can work with it as if it were created locally. This capability is fundamental to modern distributed systems. Every API call, microservice interaction, message queue, event or cache lookup depends on data being serialized before transmission and deserialized upon receive.

From an architectural perspective, serialization acts as a bridge between application memory and the outside world. Without it services could not communicate effectively, systems built in different technologies would struggle to inter operate and scalable distributed architecture simply would not be possible. Understanding this serialization deserialization cycle is the foundation for evaluating different data formats and their trade offs.


Common Serialization Formats

Once we decide to serialize the data, the next thing to do is to decide what format we should use. The answer depends on the trade-off that we are optimizing for, this could be human-readability, inter-operability or raw-performance.

→ **JSONs:** JSON has become the default choice for modern application & REST APIs, because of its simplicity, readability and it is virtually supported by every programming language. The downside is that the readability comes at a cost. JSON is text-based which means larger payload sizes and additional parsing overhead.


→ **XML:** XML was the dominant format before JSONs became popular. It is a tag-based structure which provides strong support for complex hierarchies, metadata and schema validation. However that flexibility makes XML significantly more verbose as it increases more storage and network cost. This is why today it is most commonly found in legacy enterprise systems and configuration heavy environments only.

→ **Protocol Buffers (Protobuf):** Instead of prioritising readability, protobuf is designed for efficiency. By using a compact binary format and a predefined schema, it dramatically reduces payload size and improves serialization speed. This makes it a natural fit for gRPC, micro-services communication and other high-performance distributed systems.



As architects the key lesson is that no format is universally best. JSONs optimise for simplicity, readability and inter-operability. XML for structure and validation and protobuf for performance and scalability. The choice depends on what matters most for your system.

Trade-offs: Readability vs. Efficiency vs. Compatibility




One of the most important architectural decisions when choosing a serialization format is understanding the trade off that you are making. In practice improving one characteristic often means compromising another one.

1. If developer experience and ease of debugging are your priorities, formats like JSON and XML are attractive because the data can be read directly by humans. When an API call fails or a message looks incorrect, engineers can inspect the payload immediately but the trade-off is larger message sizes and higher processing overhead.
2. If performance is the primary concern, binary formats such as protobuf and avro are far more efficient, they produce smaller payloads, consume less bandwidth and serialize much faster, which becomes increasingly valuable as traffic scales into millions of requests or events. The downside is that the data is no longer human-readable and it requires external tools to inspect.
3. The third consideration is compatibility over time. Real-world systems evolve, schema changes and multiple service versions often co-exist. Formats with strong schema support make these changes safer and much more predictable. JSON offers flexibility, but relatively weak schema enforcement while XML & protobuf provides a stronger mechanism for maintaining compatibility as your system evolves.

As architects we rarely choose a format based on a single factor, the goal is to balance readability, efficiency and long term maintainability based on specific needs of our system.

Serialization in action

Serialization becomes truly important when we look at how modern systems use it in production, almost every distributed application relies on serialization whenever data processes process, service or storage boundaries.



In the API world JSON dominates REST because of its simplicity, language independency and readability. However, system scales and performance becomes critical, many applications adopt gRPC with protobuf to reduce the payload size and improve the communication efficiency. XML still exists in an environment that depends on SOAP and other legacy enterprise integration.

The same pattern appears in caching and storage systems, before data is placed onto redis or memcache, it is typically serialized so it can be stored and retrieved efficiently. The choice of format affects both memory consumption and network overhead. Databases also leverage specialized serialization formats. For example, MongoDB uses BSON (Binary Representation of JSON) to improve storage and query performance while preserving the familiar document model. At an even larger scale, such as event driven platforms and big data pipelines, efficient serialization becomes a major cost and performance factor. Compact formats like protobuf help to reduce storage requirements, speed up data processing and support schema evolution as systems grow old over time.

So the key takeaway is that serialization is not just a data conversion technique, it is a foundational building block that directly impacts your API performance, the storage efficiency, the scalability and inter-operability across the entire system.

Summary & Key Takeaways

- Serialization enables efficient data exchange & storage.
- Choose the right format based on performance, readability, and compatibility.
- JSON for APIs, Protobuf for efficiency, Avro for big data.
- Impacts bandwidth, processing speed, and storage efficiency.
- What's next:
 - CORS: Cross-Origin Resource Sharing & Web Security

Important Interview Questions on Serialization

✚ Understanding Serialization

- What is serialization, and why is it needed in system design?
- How does serialization impact data exchange and storage?

✚ Comparing Serialization Formats

- What are the key differences between JSON, XML, Protocol Buffers, and Avro?
- When would you choose Protobuf over JSON?

✚ Serialization in APIs and Distributed Systems

- How does serialization impact API performance and efficiency?
- Why is Protocol Buffers commonly used in gRPC instead of JSON?
- How does serialization affect caching strategies in systems like Redis?

✚ Performance and Trade-offs

- What are the trade-offs between readability, efficiency, and compatibility in serialization formats?
- How does serialization impact CPU and memory usage?

✚ Real-World Applications

- How is Avro beneficial in big data systems?
- Why do some databases like MongoDB use BSON instead of JSON?

✚ Security Considerations

- What security risks are associated with serialization?
- How can improper deserialization lead to vulnerabilities?